

# CNN

May 19, 2019

```
In [1]: import os
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

import torch
import torchvision
from torchvision.datasets import ImageFolder
import torchvision.transforms as transforms

import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

```
In [9]: #Setup Torch CUDA torch device
device = torch.device('cuda:0')
```

## 0.1 Dataset

```
In [3]: transform = transforms.Compose([
    transforms.ToTensor(), # Transform to tensor
    transforms.Normalize([0.5,0.5,0.5],[0.5,0.5,0.5])
    #transforms.Normalize((0.5,),(0.5,)) # Min-max scaling to [-1, 1]
])

data_dir = os.path.join('fruits')
print('Data stored in %s' % data_dir)

trainset = ImageFolder("./fruits/Training",transform=transform)
testset = ImageFolder("./fruits/Test",transform=transform)
```

Data stored in fruits

```
In [4]: def generate_labels():
    trainset_labels = []
    testset_labels = []
    for i in trainset.imgs:
```

```

        trainset_labels.append(i[1])

    for j in testset.imgs:
        testset_labels.append(j[1])

    return (trainset_labels, testset_labels)

In [5]: # Total classes
classes_idx_dict = trainset.class_to_idx # {'Class Name': idx }
classes = len(trainset.classes)
len_trainset = len(trainset)
len_testset = len(testset)

train_labels, test_labels = generate_labels()
print(f'Trainset has total of {classes} classes')

Trainset has total of 103 classes

In [6]: trainloader = torch.utils.data.DataLoader(trainset, batch_size=60, shuffle=True)
testloader = torch.utils.data.DataLoader(testset, batch_size=60, shuffle=False)

image_shape = iter(trainloader).next()[0].shape
_, CHANNELS, HEIGHT, WIDTH = iter(trainloader).next()[0].shape
print(f'Image: batch size={image_shape[0]}, channels={image_shape[1]}, image height={ima

Image: batch size=60, channels=3, image height=100, image width=100

In [7]: class CCNet(nn.Module):
    def __init__(self):
        """
        Args:
            n_channels (int): Number of channels in the first convolutional layer. The num
                               following layers are the multipliers of n_channels.
        """
        super(CCNet, self).__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 16, kernel_size=5),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(16, 32, kernel_size=5),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(32, 64, kernel_size=5),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(64, 128, kernel_size=5),
            nn.ReLU(),

```

```

        nn.MaxPool2d(kernel_size=3, stride=3)
    )
    self.fc1 = nn.Sequential(
        nn.Linear(128, 1024),
        nn.ReLU(),
        nn.Dropout2d(p=0.8),
        nn.Linear(1024, 256),
        nn.ReLU(),
        nn.Dropout2d(p=0.8),
        nn.Linear(256, classes),
    )

    def forward(self, x, verbose=False):
        """You can (optionally) print the shapes of the intermediate variables with verb
        x = self.conv1(x)
        x = x.view(x.size(0), -1)
        x = self.fc1(x)
        return x

```

In [10]: *# Let's test the shapes of the tensors*

```

net = CCNet()
net.to(device)

with torch.no_grad():
    dataiter = iter(trainloader)
    images, labels = dataiter.next()
    images = images.to(device)
    print('Shape of the input tensor:', images.shape)

    y = net(images, verbose=True)
    print(y.shape)
    assert y.shape == torch.Size([60, classes]), f'Bad shape of y: y.shape={y.shape}'

    print('The shapes seem to be ok.')

```

Shape of the input tensor: torch.Size([60, 3, 100, 100])  
 torch.Size([60, 103])  
 The shapes seem to be ok.

In [134]: `def compute_accuracy(net, testloader):`

```

    net.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in testloader:
            images, labels = images.to(device), labels.to(device)

```

```

        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    return correct / total

In [135]: initial_learning_rate = 0.001
          final_learning_rate = 0.00001
          learning_rate = initial_learning_rate

          criterion = nn.CrossEntropyLoss()
          optimizer = optim.Adam(net.parameters(), lr=learning_rate)

In [136]: n_epochs=75
          net.train()
          for epoch in range(n_epochs):
              running_loss = 0.0
              print_every = 200  # mini-batches
              for i, (inputs, labels) in enumerate(trainloader, 0):
                  # Transfer to GPU
                  inputs, labels = inputs.to(device), labels.to(device)

                  # zero the parameter gradients
                  optimizer.zero_grad()

                  # forward + backward + optimize
                  outputs = net(inputs)
                  loss = criterion(outputs, labels)
                  loss.backward()
                  optimizer.step()

                  # print statistics
                  running_loss += loss.item()
                  if (i % print_every) == (print_every-1):
                      print('[%d, %5d] loss: %.3f' % (epoch+1, i+1, running_loss/print_every))
                      running_loss = 0.0

                  # Print accuracy after every epoch
                  accuracy = compute_accuracy(net, testloader)
                  print(f'Accuracy of the network on the {len_testset} test images: {100 * accuracy}%

          print('Finished Training')

[1, 200] loss: 4.283
[1, 400] loss: 3.528
[1, 600] loss: 2.849
[1, 800] loss: 2.326
Accuracy of the network on the 17845 test images: 53.28103110114878%

```

[2, 200] loss: 0.752  
 [2, 400] loss: 0.273  
 [2, 600] loss: 0.151  
 [2, 800] loss: 0.068  
 Accuracy of the network on the 17845 test images: 91.3981507425049%  
 [3, 200] loss: 0.075  
 [3, 400] loss: 0.082  
 [3, 600] loss: 0.047  
 [3, 800] loss: 0.027  
 Accuracy of the network on the 17845 test images: 90.92743065284394%  
 [4, 200] loss: 0.077  
 [4, 400] loss: 0.065  
 [4, 600] loss: 0.035  
 [4, 800] loss: 0.019  
 Accuracy of the network on the 17845 test images: 95.50574390585598%  
 [5, 200] loss: 0.057  
 [5, 400] loss: 0.062  
 [5, 600] loss: 0.050  
 [5, 800] loss: 0.052  
 Accuracy of the network on the 17845 test images: 93.65648641075933%  
 [6, 200] loss: 0.006  
 [6, 400] loss: 0.003  
 [6, 600] loss: 0.002  
 [6, 800] loss: 0.002  
 Accuracy of the network on the 17845 test images: 95.78033062482488%  
 [7, 200] loss: 0.001  
 [7, 400] loss: 0.000  
 [7, 600] loss: 0.035  
 [7, 800] loss: 0.161  
 Accuracy of the network on the 17845 test images: 93.19137013168954%  
 [8, 200] loss: 0.030  
 [8, 400] loss: 0.063  
 [8, 600] loss: 0.060  
 [8, 800] loss: 0.017  
 Accuracy of the network on the 17845 test images: 93.63407116839451%  
 [9, 200] loss: 0.026  
 [9, 400] loss: 0.040  
 [9, 600] loss: 0.005  
 [9, 800] loss: 0.005  
 Accuracy of the network on the 17845 test images: 93.35388063883441%  
 [10, 200] loss: 0.007  
 [10, 400] loss: 0.005  
 [10, 600] loss: 0.012  
 [10, 800] loss: 0.111  
 Accuracy of the network on the 17845 test images: 94.1552255533763%  
 [11, 200] loss: 0.068  
 [11, 400] loss: 0.019  
 [11, 600] loss: 0.003

```

[11, 800] loss: 0.015
Accuracy of the network on the 17845 test images: 91.60549173437937%
[12, 200] loss: 0.082
[12, 400] loss: 0.037
[12, 600] loss: 0.024
[12, 800] loss: 0.011
Accuracy of the network on the 17845 test images: 91.7119641356122%
[13, 200] loss: 0.049
[13, 400] loss: 0.009
[13, 600] loss: 0.011
[13, 800] loss: 0.025
Accuracy of the network on the 17845 test images: 95.79153824600728%
[14, 200] loss: 0.017
[14, 400] loss: 0.009
[14, 600] loss: 0.000
[14, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 95.86999159428412%
[15, 200] loss: 0.000
[15, 400] loss: 0.000
[15, 600] loss: 0.000
[15, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 95.93723732137853%
[16, 200] loss: 0.000
[16, 400] loss: 0.000
[16, 600] loss: 0.000
[16, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 95.95965256374335%
[17, 200] loss: 0.000
[17, 400] loss: 0.000
[17, 600] loss: 0.000
[17, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 95.98767161669936%
[18, 200] loss: 0.000
[18, 400] loss: 0.000
[18, 600] loss: 0.000
[18, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 95.99327542729056%
[19, 200] loss: 0.000
[19, 400] loss: 0.000
[19, 600] loss: 0.000
[19, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.01008685906416%
[20, 200] loss: 0.000
[20, 400] loss: 0.000
[20, 600] loss: 0.000
[20, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.00448304847296%
[21, 200] loss: 0.000

```

```

[21,  400] loss: 0.000
[21,  600] loss: 0.000
[21,  800] loss: 0.000
Accuracy of the network on the 17845 test images: 95.95965256374335%
[22,  200] loss: 0.000
[22,  400] loss: 0.000
[22,  600] loss: 0.000
[22,  800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.1053516391146%
[23,  200] loss: 0.000
[23,  400] loss: 0.000
[23,  600] loss: 0.000
[23,  800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.04370972261138%
[24,  200] loss: 0.000
[24,  400] loss: 0.000
[24,  600] loss: 0.000
[24,  800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.16138974502661%
[25,  200] loss: 0.000
[25,  400] loss: 0.000
[25,  600] loss: 0.000
[25,  800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.05491734379378%
[26,  200] loss: 0.000
[26,  400] loss: 0.000
[26,  600] loss: 0.000
[26,  800] loss: 0.000
Accuracy of the network on the 17845 test images: 95.8307649201457%
[27,  200] loss: 0.000
[27,  400] loss: 0.000
[27,  600] loss: 0.235
[27,  800] loss: 0.086
Accuracy of the network on the 17845 test images: 93.23059680582796%
[28,  200] loss: 0.022
[28,  400] loss: 0.005
[28,  600] loss: 0.008
[28,  800] loss: 0.033
Accuracy of the network on the 17845 test images: 95.57298963295041%
[29,  200] loss: 0.023
[29,  400] loss: 0.043
[29,  600] loss: 0.074
[29,  800] loss: 0.040
Accuracy of the network on the 17845 test images: 94.4578313253012%
[30,  200] loss: 0.029
[30,  400] loss: 0.006
[30,  600] loss: 0.018
[30,  800] loss: 0.034

```

Accuracy of the network on the 17845 test images: 92.86634911739982%  
 [31, 200] loss: 0.016  
 [31, 400] loss: 0.012  
 [31, 600] loss: 0.006  
 [31, 800] loss: 0.082  
 Accuracy of the network on the 17845 test images: 94.49145418884841%  
 [32, 200] loss: 0.012  
 [32, 400] loss: 0.003  
 [32, 600] loss: 0.005  
 [32, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.61782011768003%  
 [33, 200] loss: 0.000  
 [33, 400] loss: 0.000  
 [33, 600] loss: 0.000  
 [33, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.85878397310171%  
 [34, 200] loss: 0.000  
 [34, 400] loss: 0.000  
 [34, 600] loss: 0.000  
 [34, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.90921826842252%  
 [35, 200] loss: 0.000  
 [35, 400] loss: 0.000  
 [35, 600] loss: 0.000  
 [35, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.99887923788177%  
 [36, 200] loss: 0.000  
 [36, 400] loss: 0.000  
 [36, 600] loss: 0.000  
 [36, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 96.02689829083776%  
 [37, 200] loss: 0.000  
 [37, 400] loss: 0.000  
 [37, 600] loss: 0.000  
 [37, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 96.06612496497618%  
 [38, 200] loss: 0.000  
 [38, 400] loss: 0.000  
 [38, 600] loss: 0.000  
 [38, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 96.16138974502661%  
 [39, 200] loss: 0.000  
 [39, 400] loss: 0.000  
 [39, 600] loss: 0.000  
 [39, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 96.17820117680023%  
 [40, 200] loss: 0.000  
 [40, 400] loss: 0.000



```

[40, 600] loss: 0.000
[40, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.17259736620902%
[41, 200] loss: 0.000
[41, 400] loss: 0.000
[41, 600] loss: 0.000
[41, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.27346595685066%
[42, 200] loss: 0.000
[42, 400] loss: 0.000
[42, 600] loss: 0.000
[42, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.26225833566825%
[43, 200] loss: 0.000
[43, 400] loss: 0.000
[43, 600] loss: 0.000
[43, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.28467357803306%
[44, 200] loss: 0.000
[44, 400] loss: 0.000
[44, 600] loss: 0.000
[44, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.3687307369011%
[45, 200] loss: 0.238
[45, 400] loss: 0.085
[45, 600] loss: 0.021
[45, 800] loss: 0.030
Accuracy of the network on the 17845 test images: 95.34883720930233%
[46, 200] loss: 0.027
[46, 400] loss: 0.039
[46, 600] loss: 0.006
[46, 800] loss: 0.001
Accuracy of the network on the 17845 test images: 94.28971700756514%
[47, 200] loss: 0.057
[47, 400] loss: 0.011
[47, 600] loss: 0.008
[47, 800] loss: 0.008
Accuracy of the network on the 17845 test images: 96.49201456990754%
[48, 200] loss: 0.005
[48, 400] loss: 0.006
[48, 600] loss: 0.054
[48, 800] loss: 0.084
Accuracy of the network on the 17845 test images: 91.43177360605212%
[49, 200] loss: 0.050
[49, 400] loss: 0.028
[49, 600] loss: 0.001
[49, 800] loss: 0.001
Accuracy of the network on the 17845 test images: 95.42729055757916%

```

```

[50, 200] loss: 0.000
[50, 400] loss: 0.001
[50, 600] loss: 0.000
[50, 800] loss: 0.055
Accuracy of the network on the 17845 test images: 90.93863827402635%
[51, 200] loss: 0.125
[51, 400] loss: 0.055
[51, 600] loss: 0.073
[51, 800] loss: 0.011
Accuracy of the network on the 17845 test images: 96.14457831325302%
[52, 200] loss: 0.000
[52, 400] loss: 0.000
[52, 600] loss: 0.000
[52, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.81143177360605%
[53, 200] loss: 0.000
[53, 400] loss: 0.000
[53, 600] loss: 0.000
[53, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.89548893247408%
[54, 200] loss: 0.000
[54, 400] loss: 0.000
[54, 600] loss: 0.000
[54, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.87307369010928%
[55, 200] loss: 0.000
[55, 400] loss: 0.070
[55, 600] loss: 0.188
[55, 800] loss: 0.045
Accuracy of the network on the 17845 test images: 93.0456710563183%
[56, 200] loss: 0.031
[56, 400] loss: 0.006
[56, 600] loss: 0.008
[56, 800] loss: 0.009
Accuracy of the network on the 17845 test images: 93.44354160829363%
[57, 200] loss: 0.051
[57, 400] loss: 0.023
[57, 600] loss: 0.003
[57, 800] loss: 0.020
Accuracy of the network on the 17845 test images: 94.08237601569067%
[58, 200] loss: 0.026
[58, 400] loss: 0.069
[58, 600] loss: 0.031
[58, 800] loss: 0.052
Accuracy of the network on the 17845 test images: 94.04314934155225%
[59, 200] loss: 0.022
[59, 400] loss: 0.052
[59, 600] loss: 0.006

```

[59, 800] loss: 0.028  
 Accuracy of the network on the 17845 test images: 94.07677220509947%  
 [60, 200] loss: 0.005  
 [60, 400] loss: 0.002  
 [60, 600] loss: 0.008  
 [60, 800] loss: 0.015  
 Accuracy of the network on the 17845 test images: 92.86074530680864%  
 [61, 200] loss: 0.078  
 [61, 400] loss: 0.047  
 [61, 600] loss: 0.064  
 [61, 800] loss: 0.005  
 Accuracy of the network on the 17845 test images: 95.09666573269824%  
 [62, 200] loss: 0.008  
 [62, 400] loss: 0.016  
 [62, 600] loss: 0.006  
 [62, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.44410198935276%  
 [63, 200] loss: 0.000  
 [63, 400] loss: 0.000  
 [63, 600] loss: 0.000  
 [63, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.64583917063604%  
 [64, 200] loss: 0.000  
 [64, 400] loss: 0.000  
 [64, 600] loss: 0.000  
 [64, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.8475763519193%  
 [65, 200] loss: 0.000  
 [65, 400] loss: 0.000  
 [65, 600] loss: 0.000  
 [65, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 95.98206780610815%  
 [66, 200] loss: 0.000  
 [66, 400] loss: 0.000  
 [66, 600] loss: 0.000  
 [66, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 96.03250210142897%  
 [67, 200] loss: 0.000  
 [67, 400] loss: 0.000  
 [67, 600] loss: 0.000  
 [67, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 96.07733258615859%  
 [68, 200] loss: 0.000  
 [68, 400] loss: 0.000  
 [68, 600] loss: 0.000  
 [68, 800] loss: 0.000  
 Accuracy of the network on the 17845 test images: 96.0941440179322%  
 [69, 200] loss: 0.000

```

[69, 400] loss: 0.000
[69, 600] loss: 0.000
[69, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.13337069207061%
[70, 200] loss: 0.000
[70, 400] loss: 0.000
[70, 600] loss: 0.000
[70, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.13337069207061%
[71, 200] loss: 0.000
[71, 400] loss: 0.000
[71, 600] loss: 0.000
[71, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.16138974502661%
[72, 200] loss: 0.000
[72, 400] loss: 0.000
[72, 600] loss: 0.000
[72, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.16699355561782%
[73, 200] loss: 0.000
[73, 400] loss: 0.000
[73, 600] loss: 0.000
[73, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.20622022975624%
[74, 200] loss: 0.000
[74, 400] loss: 0.000
[74, 600] loss: 0.000
[74, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.20061641916503%
[75, 200] loss: 0.000
[75, 400] loss: 0.000
[75, 600] loss: 0.000
[75, 800] loss: 0.000
Accuracy of the network on the 17845 test images: 96.27346595685066%
Finished Training

```

```

In [137]: accuracy = compute_accuracy(net, testloader)
          print('Accuracy of the network on the test images: %.3f' % accuracy)

```

```

Accuracy of the network on the test images: 0.963

```

```

In [139]: filename = 'cnn.pth'

          try:
              do_save = input('Do you want to save the model (type yes to confirm)? ').lower()
              if do_save == 'yes':

```

```

        torch.save(net.state_dict(), filename)
        print('Model saved to %s' % filename)
    else:
        print('Model not saved')
except:
    raise Exception('The notebook should be run or validated with skip_training=True.')

```

Do you want to save the model (type yes to confirm)? yes  
Model saved to cnn.pth

```

In [ ]: net = CCNet()
        net.load_state_dict(torch.load(filename, map_location=lambda storage, loc: storage))
        net.to(device)
        print(f'Model loaded from {filename}')

```